Type-Directed Program Synthesis and Constraint Generation for Accelerator Library Portability

Bruce Collie
School of Informatics
University of Edinburgh
United Kingdom
bruce.collie@ed.ac.uk

Philip Ginsbach
School of Informatics
University of Edinburgh
United Kingdom
philip.ginsbach@ed.ac.uk

Michael F.P. O'Boyle
School of Informatics
University of Edinburgh
United Kingdom
mob@inf.ed.ac.uk

Abstract—**Fast numerical libraries have been a cornerstone of scientific computing for decades, but this comes at a price. Programs may be tied to vendor specific software ecosystems resulting in polluted, non-portable code. As we enter an era of heterogeneous computing, there is an explosion in the number of accelerator libraries required to harness specialized hardware. We need a system that allows developers to exploit ever-changing accelerator libraries, without over-specializing their code.**

As we cannot know the behavior of future libraries ahead of time, this paper develops a scheme that assists developers in matching their code to new libraries, without requiring the source code for these libraries.

Furthermore, it can recover equivalent code from programs that use existing libraries and automatically port them to new interfaces. It first uses program synthesis to determine the meaning of a library, then maps the synthesized description into generalized constraints which are used to search the program for replacement opportunities to present to the developer.

We applied this approach to existing large applications from the scientific computing and deep learning domains. Using our approach, we show speedups ranging from $1.1\times$ to over $10\times$ on end to end performance when using accelerator libraries.

Index Terms—**program synthesis, code rejuvenation, constraint programming, compilers**

I. Introduction

Fast numerical libraries have been a cornerstone of scientific computing for decades [1], [2]. They provide efficient implementations of key algorithmic components and allow a separation of concerns. This comes at a cost, however, as it may tie programs into vendor-specific software ecosystems and results in non-portable, polluted code. A new library API means that the original "vanilla" code has to be recovered, and then modified to use the new libraries.

The risk of being tied into an out of date library API has led some developers to release multiple versions of their code, e.g. PyTorch [3] and Darknet [4]. This requires the maintenance of multiple code bases and complex build systems. However, as we witness the rise of specialized heterogeneous accelerators, we also see a proliferation of accelerator libraries [5], [6], [7], [8]. In the long-term, a multi-versioned code base is not sustainable.

This paper develops an alternative approach: a compiler-based scheme that discovers opportunities to use new accelerator libraries in user code, with little prior knowledge of the libraries. Furthermore, it can recover behaviorally equivalent code from programs that use existing libraries and automatically port them to new interfaces. In order to reduce developer burden, it attempts to do this with minimal intervention using program synthesis and graph matching.

Program synthesis is a well studied area that deals with searching a program space to find candidate programs that match a specification [9]. Our program synthesis implementation uses a number of generic control-flow components and a set of heuristics defining when they should be applied. These heuristics are driven by a library's type signature and lightweight annotations provided by the library vendor. Crucially, these annotations are easily extracted from documentation and require no knowledge of a library's internals.

Once we know what a library does, we need to see if the developer's program has structures that match its behavior. We achieve this by automatically describing the synthesized program as a set of constraints which we then use to search the user code. As the synthesized program may not easily match existing code, we first generate many equivalent versions, normalize and then generalize them to a common description that determines the most appropriate constraints. When we find a match we suggest to the developer that a replacement could be made to take advantage of a different library.

One unique aspect of our approach is that synthesis and constraint generation is entirely at the LLVM level, which allows us to target large existing code bases and show significant performance improvement. We applied this approach to existing large applications from scientific computing and deep learning written in C, C++ and Fortran. We show speedups ranging from $1.1\times$ to over $10\times$ improvement when implementing replacements suggested by our tools.

II. Example

This section illustrates how our approach helps in porting code that uses an existing library API to a new library API with increased functionality. Consider the code sample on the left of Figure 1. This is an inner loop taken from a subroutine in NWChem [10], a widely used chemical simulation suite

```
Original: BLAS

do i=1,nbeads2
  t = (i-1)/dble(nbeads2-1)

  j1 = t*(nbeads1-1) + 1
  j2 = j1+1
  t1 = (j1-1)/dble(nbeads1-1)
  t2 = (j2-1)/dble(nbeads1-1)
  t3 = (t-t1)/(t2-t1)

  if (j2.gt.nbeads1) then
      t3 = 0.0d0
      j2=nbeads1
  end if

  shift = (j1-1)*3*nion
  call dcopy(3*nion,dbl_mb(c(1)+shift),
          1,dbl_mb(r1(1)),1)
  call dcopy(3*nion,dbl_mb(r1(1)),
          1,dbl_mb(r3(1)),1)
  do j=0,3*nion-1
    dbl_mb(r3(1)+j) = (1.0d0-t3) *
      dbl_mb(r3(1)+j)
  end do
  shift = (j2-1)*3*nion
  call dcopy(3*nion,dbl_mb(c(1)+shift),
          1,dbl_mb(r2(1)),1)
  call daxpy(3*nion,t3,dbl_mb(r2(1)),
          1,dbl_mb(r3(1)),1)
end do
```

```
SYNTHESIZER

MODEL FOR:
mkl_daxpby(n,a,x,incx,b,y,incy)

MODEL FOR:
daxpy(n,a,x,incx,y,incy)
```

UNCHANGED

```
do j=0,3*nion-1,1
  dbl_mb(r3(1)+j) = (1.0d0-t3) *
    dbl_mb(r3(1)+j)
end do
...
do j=0,3*nion-1,1
  dbl_mb(r3(1)+j) =
    t3*dbl_mb(r2(1)+j) +
    dbl_mb(r3(1)+j)
end do
```

INLINE

PATTERN MATCH + REPLACE CODE

```
Transformed: MKL

do i=1,nbeads2
  t = (i-1)/dble(nbeads2-1)

  j1 = t*(nbeads1-1) + 1
  j2 = j1+1
  t1 = (j1-1)/dble(nbeads1-1)
  t2 = (j2-1)/dble(nbeads1-1)
  t3 = (t-t1)/(t2-t1)

  if (j2.gt.nbeads1) then
      t3 = 0.0d0
      j2=nbeads1
  end if

  shift = (j1-1)*3*nion
  call dcopy(3*nion,dbl_mb(c(1)+shift),
          1,dbl_mb(r1(1)),1)

  shift = (j2-1)*3*nion
  call dcopy(3*nion,dbl_mb(c(1)+shift),
          1,dbl_mb(r2(1)),1)
  call dcopy(3*nion,dbl_mb(r1(1)),
          1,dbl_mb(r3(1)),1)

  call mkl_daxpby(3*nion,
          t3,dbl_mb(r2(1)),1,
          (1.0d0-t3),dbl_mb(r3(1)),1)
end do
```

Fig. 1: Porting BLAS to MKL, managing API evolution. On the left is code taken from NWChem [10], a widely-used chemical simulation suite. Our approach learns and inlines the code for `daxpy`, then identifies the resulting code as being equivalent to the model learned for `mkl_daxpby`.

that makes explicit calls to BLAS libraries. The code contains manual loops over arrays and calls to BLAS routines (`dcopy` and `daxpy`). We wish to port this code to use Intel's MKL libraries (as shown on the right hand side, which makes use of the extended MKL functionality `mkl_daxpby`).

On the left hand side of the figure, there are two highlighted sections of code. The first highlighted piece of code is a simple loop that performs the following abstract vector operation:

$$\mathbf{r_3} \leftarrow (1 - t_3)\mathbf{r_3} \tag{1}$$

The second highlighted piece of code is a call to `daxpy`. *If we had access to the `daxpy` source code, we would see this corresponds to the following vector operation:*

$$\mathbf{r_3} \leftarrow t_3\mathbf{r_2} + \mathbf{r_3} \tag{2}$$

As we wish to port NWChem to MKL, we can exploit the extended MKL [11] BLAS function `mkl_daxpby`. This supports vector scaling, combined with `daxpy` corresponding to the vector operation

$$\mathbf{r} \leftarrow (1 - t)\mathbf{r} + t\mathbf{r}' \tag{3}$$

A call to this extended implementation `mkl_daxpby` is shown in the highlighted box on the RHS of Figure 1. The two original operations in equations (1,2) can be rewritten as a single operation by substituting for $\mathbf{r_3}$ in equation (2) to give

$$\mathbf{r_3} \leftarrow t_3\mathbf{r_2} + (1 - t_3)\mathbf{r_3} \tag{4}$$

which corresponds to the extended function `daxpby`. In this example, it means that we can legally replace the two

highlighted pieces of code on the LHS of the figure with the similarly highlighted code on the RHS.

Match and Replace: If there is a source level description of both versions of `daxpy`, we first inline the original call, as shown in the **Inline** box in Figure 1 to give the modified code in the central part of the figure.[1] We then try to match this modified code to the code corresponding to the extended version from MKL [11]. We achieve this using a novel graph based constraint solver which matches the two codes and suggests to the developer that the old calls be replaced with the new one.

Synthesis: In practice, we cannot guarantee there is a suitable source level description of every use of a library. This may be due to the library provider not releasing an appropriate description; it no longer being available; or being poorly documented. It may also be defined in a manner suitable for human consumption but not compiler automation. It is certainly the case that there is not agreement among all library developers about a universal language to describe the semantics of all their libraries.

If the source code of the two libraries is unavailable, we use program synthesis to generate programs that correspond to both versions of `daxpy` as shown by the **Synthesizer** box in Figure 1. We then, as before, inline the original BLAS `daxpy`, pattern match and suggest to the developer that it be replaced with a call to the extended MKL version.

---

[1] Calls to other functions (e.g. `dcopy`) will also be inlined, but are not shown in the figure for clarity.

Library Centric

LIBRARY FUNCTIONS

```
dgemv(...)
opt_dgemv(...)
softmax(...)
fir_1d(...)
relu(...)
dot(...)
```

→ **SYNTHESIZE** →

SYNTHESIZED IMPLEMENTATIONS

```
declare void dgemv(...) {
  //synthesized code
}
...
declare void dot(...) {
  //sythesized code
}
```

→ **GENERALIZE** →

CONSTRAINTS

```
Constraint dgemv (
  {x} is add inst
  and ...
)
Constraint dot (
  {x} is mul inst
  and ...
)
```

INLINE ↓ (from SYNTHESIZED IMPLEMENTATIONS)

SEARCH ↓ (from CONSTRAINTS)

Program Centric

USER CODE

```
do i=1,nlocs
  val = val+x(i)*y(i)
end do
call dgemv(..., x, ...)
```

→ **INLINE** →

INLINED USER CODE

```
do i=1,nlocs
  val = val+x(i)*y(i)
end do
do i=1,N
  do j=1,M
    // dgemv loop
    ...
  end do
end do
```

→ **SEARCH** →

MATCH RESULTS

```
found "dot" at ...
  with (x, y)
...
found "dgemv" at ...
  with ...
```

REPLACE ↓

IMPROVED USER CODE

```
call dot(nlocs, x, y)
call opt_dgemv(..., x, ...)
```

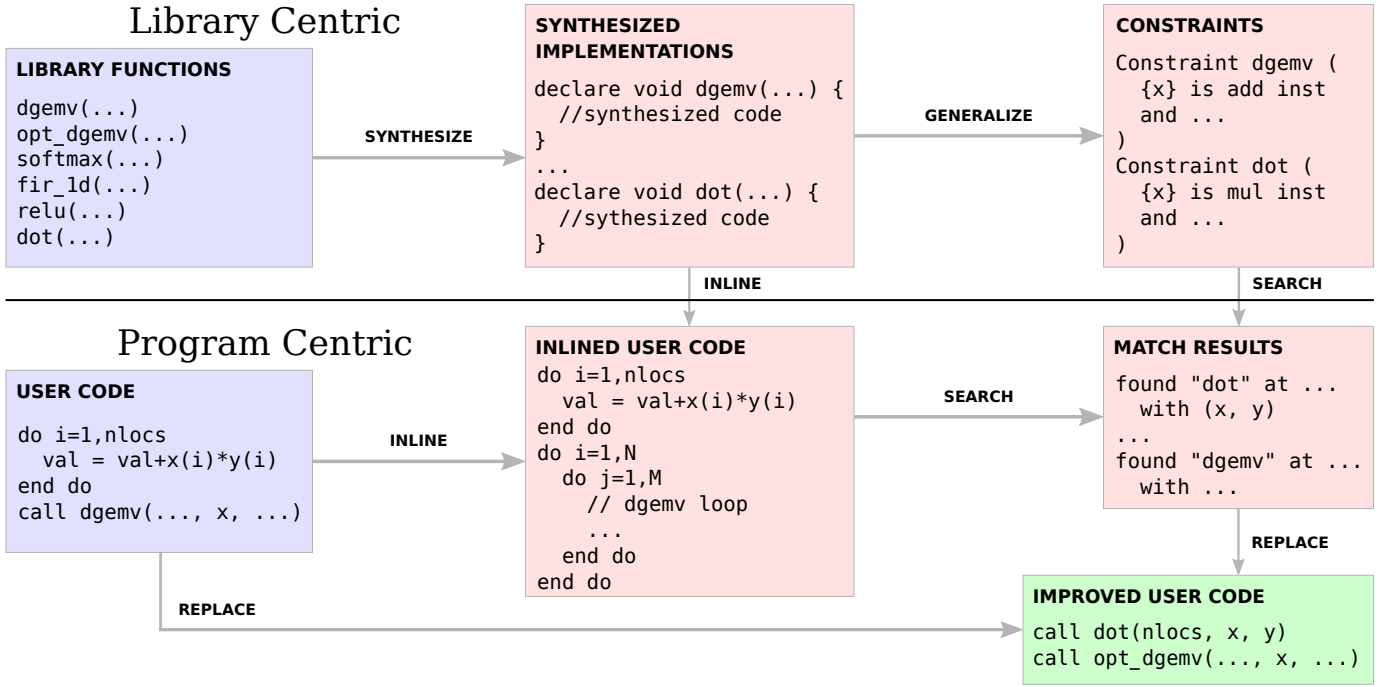**REPLACE** (from USER CODE to IMPROVED USER CODE)

Fig. 2: An overview of the data flow through the different stages of our system. Inputs are shown in blue on the left, intermediate products in red and the final product in green in the bottom right corner. For brevity, the content shown in each box is an approximation of the actual data in the system.

Thus, the developer is able to port their code to a new, extended library without having to identify the opportunity manually (they need only agree to the suggested replacement). In this example, it results in a 20% performance improvement on an Intel Xeon E5-2620. If the code is to be ported again or an improved library is released, then the procedure can be repeated, avoiding legacy API tie-in. At the heart of our approach is the use of program synthesis and graph based generalized constraint matching.

III. OVERVIEW

Figure 2 gives a high level overview of our approach. It can be split into two sections: *library centric* and *program centric*. The library centric work corresponds to the top three boxes in the diagram and has to be performed once for each new library considered. The program centric work corresponds to the remaining boxes and is performed when a user program has to be re-targeted to a new API.

A. Library Centric

Given the extended type of a library, we use oracle guided program synthesis [12] to generate a program that is equivalent to the behavior of the library. This is achieved by generating many input/output pairs which are used to guide synthesis. As the space of programs is unbounded, we exploit type information and heuristics to generate synthesized programs in a reasonable time.

Once we have a candidate synthesized implementation, it is unlikely to have an identical structure to all user code.

We therefore generate many versions and then generalize and derive a constraint program that describes the generalized program structure. These constraints will be used later to search the user programs for equivalent matches. As we work within the LLVM compiler infrastructure, we synthesize programs and derive constraints at the LLVM static single assignment intermediate representation of a program.

B. Program Centric

Once a user program is to be ported, we first inline the synthesized implementation of any library calls. If there is a source code description of the library available, this can be used, otherwise we rely on the synthesized code. Next we search for code patterns corresponding to the constraints derived from the library centric phase.

C. Contribution

The program centric phase is based on prior work and uses the constraint language IDL, the SMT solver and replacement technique described in [13]. Our contribution is restricted to the library centric phases where we make two novel contributions:

- Oracle-guided LLVM program synthesis of black-box imperative libraries with nested control structures using lightweight type signature annotations.
- Automatic generation and generalization of program constraints from examples using graph matching.

We elaborate on these in Section IV and Section V before evaluating our approach on a range of real-world applications.

D. Practical Usage

Neither the synthesis nor the generalization procedures in Figure 2 can guarantee semantic correctness; instead we rely on the notion of behavioral equivalence as described in Section IV-E. In practice, we rely on the developer to sign off on any code replacement. To avoid wasting developer time with false positives, each potential replacement can be first checked by comparing the output of the original code against the suggested replacement. Only successful candidates are then presented. Section VII describes the ways in which both the library and program centric components may exhibit unsoundness, and show that in practice the usefulness of our tools is not greatly affected.

IV. LEARNING PROGRAMS

A. Annotated Signatures

Automatically learning the behavior of an arbitrary function, given only its type signature, is generally an intractable problem. We define a simple language of annotations (in the spirit of a minimal logic programming language) that can be used by library vendors to annotate their functions with arbitrary additional semantic information not expressible in a type signature.

A useful motivating example is encoding the relationship between a pointer to allocated memory and the size of that memory. The upper right corner of Figure 3 shows this being used to annotate the function `daxpy` from Figure 1. A full listing of the annotations used in this paper is given in Section IV-B—the properties corresponding to these annotations are conceptually simple and can be easily extracted from API documentation.

B. Annotation Details

The descriptions and algorithms given above are abstract and could be used to instantiate many different program synthesizers, depending on the annotations and fragment templates used. The functions evaluated in this paper are synthesized using five core annotations, each of which is listed and briefly specified below:

size(xs, n): the pointer `xs` points to allocated memory with `n` elements.

output(x): the pointer `x` is an output parameter for the function.

enum(x, c0, ..., cN): the parameter `x` must take one of the distinct constant values `c0...cN`.

pack(xs, c): each logical entry in the array pointed to by `xs` contains `c` physical elements.

indices(xs): elements of `xs` in memory are logically array indices.

Almost all of the functions we synthesize use at least one `size` annotation. The use of separate pointer and size arguments is endemic to C function signatures, and we found documentation highlighting this relationship in every library considered. Similarly, logical output arguments are always highlighted in documentation (some degree of inference could
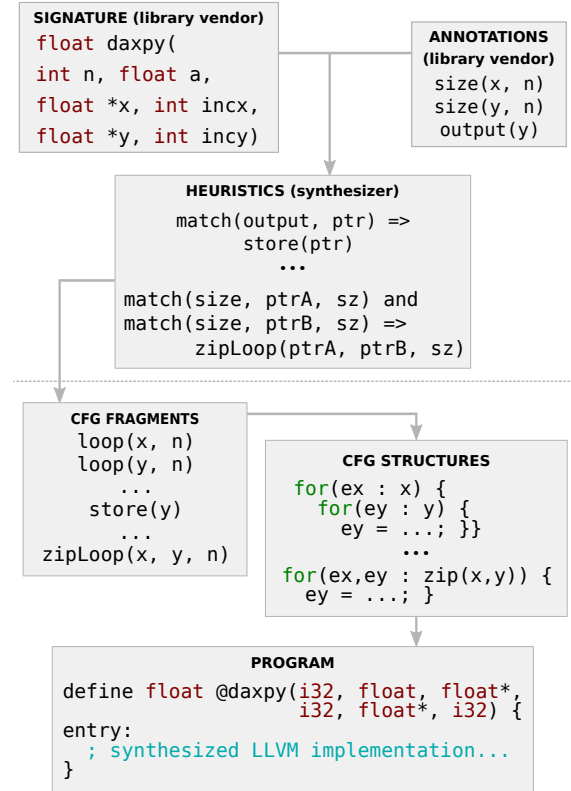

Fig. 3: A simplified illustration of how our synthesizer learns an implementation for the `daxpy` linear algebra function. Inputs to the process (from the library vendor and the synthesizer itself) are given above the dotted line, and the synthesis process is given below. The set of annotations given (`size` and `output`) is complete, but the heuristics, fragments and instantiated CFG compositions have been abbreviated for brevity. A full explanation of this worked example is given in Section IV-C

be performed for this annotation by considering **const**-specified pointers). The annotation `enum` is used for BLAS functions that perform (for example) transposed or non-transposed versions of the same computation, and we found the values were listed prominently in documentation.

The remaining two annotations are less closely tied to documentation, but were used in only a small number of cases. `pack` was used to simplify implementation details for functions dealing with homegeneous structure types, and all uses of it could easily be removed with no conceptual changes. Only one use of `indices` was necessary—when synthesizing `spmv`, to ensure that indirect memory accesses were safe. This annotation could be found in documentation, albeit less prominently than `size` and `output`.

These annotations were sufficient to synthesize all the functions described in Section VII. We do not believe supplying them represents a significant burden on the library vendor, and could easily be automated in the simpler cases.

Algorithm 1 Dataflow generation

```
 1: function FILLDATAFLOW(cfg, n)
 2:     tree ← dominance tree of cfg
 3:     phis ← ∅
 4:     for each block in inorder(tree) do
 5:         if block has > 1 predecessors then
 6:             phis ← phis ∪ { new φ node in block }
 7:         end if
 8:         if block is a dataflow block then
 9:             live ← live SSA values in block
10:             generate n instructions sampling from live
11:         end if
12:     end for
13:     for each φ in phis do
14:         live ← live SSA values at block with φ
15:         choose incoming values to φ from live
16:     end for
17: end function
```

C. Control Flow

Our approach to synthesizing candidate programs is two-phase. Firstly, we construct and sample from the set of potential control-flow structures that a candidate might use. Secondly, we fill in the control structures with instructions.

a) Control Flow Fragments: A control flow fragment is a region of code with 'holes' in it, such that they can be composed with other fragments to form a complete program CFG. The holes in a fragment may be filled by any other fragment, but may also remain empty; a valid LLVM IR programs can be extracted from any composition. Additionally, fragments may be parameterized.

b) Querying Properties: In order to generate potentially valid control flow structures, we require a set of candidate fragments that might comprise a solution. This part of the synthesis process is driven by inference rules expressing heuristics on when each type of control flow fragment should be instantiated. For example, the first line in the HEURISTICS box in Figure 3 shows a rule for instantiating a loop over memory if the size of that memory is known. Queries use a limited form of unification with conjunctive matches (e.g. `sz` in Figure 3 when instantiating a `zipLoop`).

c) Sampling Control Flow: The first step in the synthesis process is to enumerate all the possible query matches against the function type signature and property annotations. This yields a set of control flow fragments, from which the synthesizer will construct program fragments by composition. To do this, we perform an exhaustive search over the possible compositions of up to 3 fragments. If no solution is found with 3 fragments, we revert to a random sampling process, although we have not found this to be necessary in practice.

d) Example: Figure 3 shows a worked example of this synthesis procedure for the `daxpy` function used in Section II. The top two boxes show its type signature and the full set of easily-obtainable library vendor annotations required for synthesis. These annotations are matched against the synthesizer's set of heuristics, shown in the central box. Only the two most important ones are shown: creating a store instruction if a pointer is a logical output, and iterating over pointers with the same logical size together. Then, the synthesizer collects the full set of possible fragments; some less useful ones such as individual loops over `x` and `y` are created at this stage. These fragments are then composed to form program structures (some examples are shown as pseudo-C to the right), and finally compiled to LLVM IR where instructions are added and testing is performed (bottom box).

D. Generating Programs

Given a control flow structure composed from several fragments, the final step in our synthesis process is to add instructions to it, producing a candidate program. We use a generic algorithm to do this, with no knowledge of the specific fragments that comprise the control flow structure. We walk the dominance tree of the control flow graph, inserting stochastically sampled instructions in the appropriate places. Additionally, φ nodes are used to handle looping or divergent control correctly. This algorithm is shown in Algorithm 1.

At each node in the dominance tree, a set of 'live' instructions is maintained. The initial instructions in this set come from fragment specifications (for example, values loaded from memory in a loop iteration). New instructions are sampled from a set of possibilities: we allow for integer and floating point arithmetic, calls to intrinsic mathematical functions, and a small number of other simple primitives such as conditional selects. Searching for a correct candidate program amounts to iteratively performing this instruction generation algorithm on each possible control flow structure, testing each resulting program for behavioral equivalence until a solution is found.

By exploiting annotated type signatures and general heuristics, we can realistically synthesize complex imperative programs. Section VII-D discusses the complexity and time requirements of this process.

E. Behavioral Equivalence

Our methodology is concerned with 'black-box' interfaces. We are able to observe input-output behavior, but have no ground truth to evaluate the correctness of potential solutions against. We judge a solution as behaviorally equivalent to a reference implementation, if it behaves equivalently across a large set of random example inputs. More formally, for a solution s, reference r and input vectors $\mathbf{x}_i$.

$$correct(s) \iff \forall i \,.\, s(\mathbf{x}_i) \approx (\mathbf{x}_i)$$

An approximate notion of equality ($\approx$) is used to compare floating point values. This correctness decision is unsound; there is no way to establish formally that a candidate program will behave correctly on every possible input. However, work on property-based testing [14] uses a similar assumption to ours—if enough is known about the way inputs are used by a program, then equivalence over a large number of random samples strongly implies equality. In practice, users could be asked to sign off on synthesized programs.
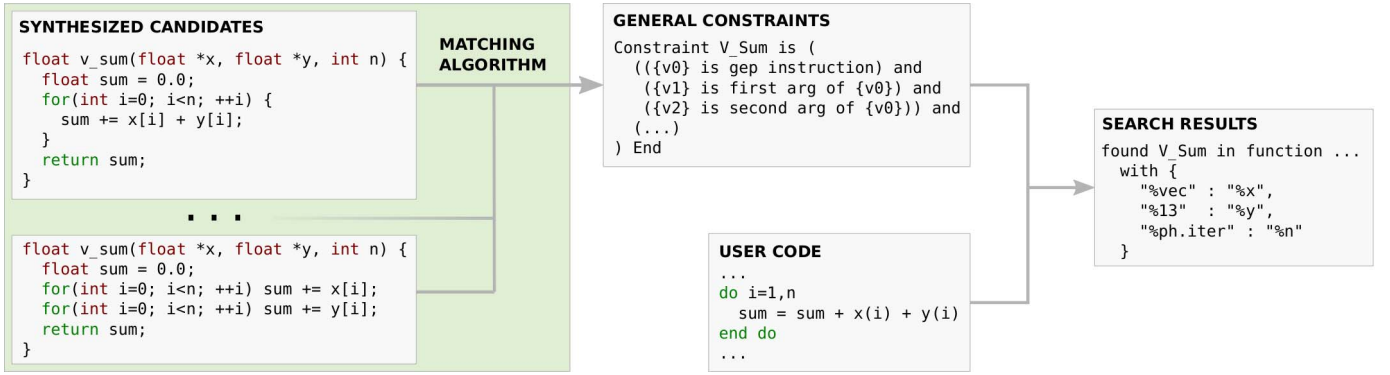
Fig. 4: Generating constraints from synthesized candidates: Our matching algorithm transforms synthesized programs into constraints and then constructs a generalized constraint description in IDL [13]. This searches user code and generates matches. Our contribution (generating constraints from examples) is highlighted in green.

V. RECOGNIZING LEARNED PROGRAMS

In Section IV we described a technique for synthesizing LLVM programs based on input-output examples and information provided by library vendors about their function call interfaces. This allows us to model existing and new target library interfaces.

We use these generated LLVM programs to develop a constraint description of the underlying library implementation. These constraints are then used to automatically detect relevant sections in existing code bases. Searching for exact matches of synthesized LLVM IR fragments is unlikely to achieve any success. The synthesized code itself is not a sufficient model, as the IR is not normalized and there are many potential alternative implementations. In order to generate a useful model, we need to extract common features from multiple implementations.

In this section we will establish an approach to converting sets of synthesized IR fragments into constraint programs as shown in Figure 4. The contribution of this paper is highlighted in the green section of the figure: synthesized programs are merged with a graph matching algorithm and a constraint description is generated from the result. This constraint description can then be applied to user code in order to detect matching instances.

Firstly, we will show how individual fragments can be converted into candidates. Secondly, we show how multiple fragments can be accumulated, crystallizing out the most important features using graph constraint matching. Thirdly, we will show how these accumulated fragments can be converted into constraints which are used to search the developer's program.

A. Generating Constraints From LLVM IR Fragments

To build a constraint description of an LLVM IR fragment, we abstract the SSA-form code into a graph $G = (V, E)$ with edges $E \subset \mathbb{N} \times V \times V$.

We use the notation $a \xrightarrow{n} b$ to signify $(n, a, b) \in E$, which means that instruction a is the nth argument of instruction b. Function parameter and constants are modeled as instructions

with special opcodes. We can now mechanically generate a constraint description that specifies the exact combination of instructions and argument relations from the program. In this specific example, the output is in the IDL programming language [13] (GENERAL CONSTRAINTS in Figure 5).

This results in a constraint program that detects exact matches. However, we want to capture a wider class of possible functionally equivalent implementations rather than a single exact program. In the next section we describe a graph-based method for generalizing these constraints.

B. Matching LLVM IR Fragments Together

In order to generate a common constraint description of a representative set of LLVM IR functions, we need to match together the respective graph representations. The aim of this is to uncover the essential parts of the computation. These are then extracted and turned into constraints.

Given a set of LLVM IR programs, we construct a single graph by directly combining their respective vertices and edges. We now aim to match together nodes from the different programs that are semantically equivalent. For this, we will derive an equivalence relation $\sim$ that partitions the vertices into a set of equivalence classes $V/\sim$. Vertices that correspond to the same behavior across examples will belong to the same class. We use the usual notation for the equivalence class $\bar{u}$ of a vertex u: $\bar{u} := \{v \in V \mid u \sim v\}$, and write $\bar{x} \xrightarrow{n} \bar{y}$ to express the following properties:

$$\forall a \in \bar{x} . \exists b \in \bar{y} . a \xrightarrow{n} b$$
$$\forall b \in \bar{x} . \exists a \in \bar{y} . a \xrightarrow{n} b$$

These properties ensure that each instruction in class $\bar{x}$ is the nth argument of an instruction in $\bar{y}$, and for each instruction in $\bar{y}$, the nth argument is in $\bar{x}$.

1) Deriving An Optimization Target: We intuitively require that an approximation of several properties should hold. Firstly, instructions that are mapped onto each other should have the same opcode:
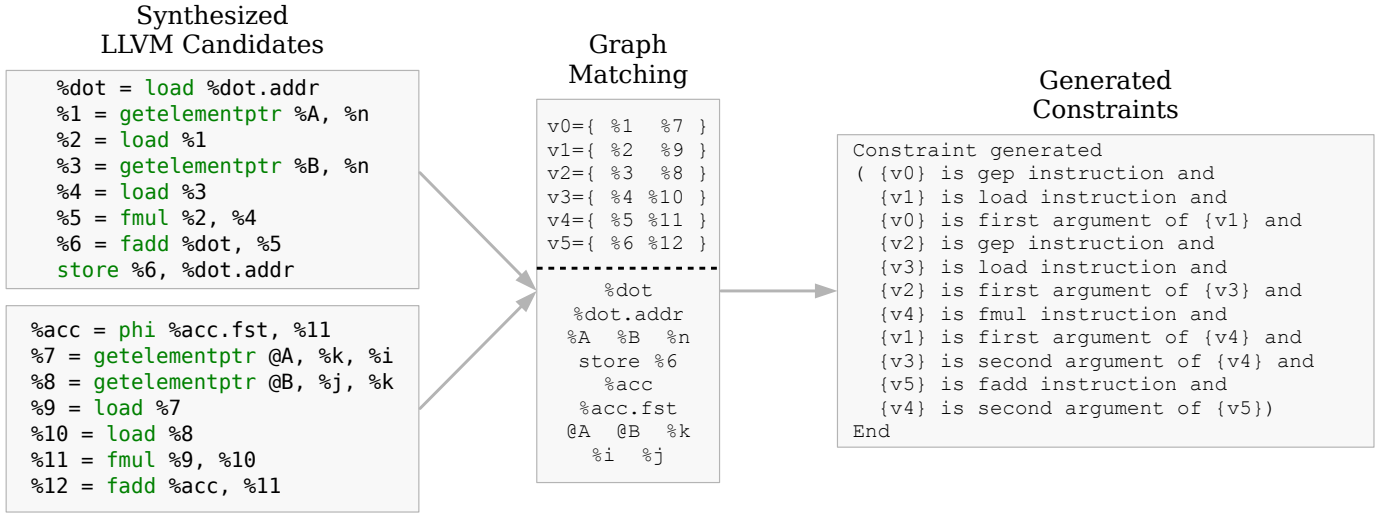
$$a \sim b \implies op(a) = op(b)$$

Fig. 5: Generalizing Constraints: Candidate constraints from two LLVM IR versions of a dot product are matched. The algorithm identifies six groups of instructions to be matched together as shown in the top half of the middle graph matching box. The others are discarded as shown in the lower half. The generalized constraints are then output as shown in the final box.

Secondly, edge relationships should be compatible with the equivalence relation:

$$x \xrightarrow{n} y \implies \bar{x} \xrightarrow{n} \bar{y}$$

Thirdly, we do not want argument relationships to be collapsed together by the equivalence relation:

$$a \xrightarrow{n} b \implies \bar{a} \neq \bar{b}$$
$$a \xrightarrow{n} c \land b \xrightarrow{m} c \land \bar{a} = \bar{b} \implies a = b \lor n = m$$

Of course, we cannot expect all the criteria to be perfectly fulfilled. Instead we aim for an approximate result. We thus define a metric m on the set of possible equivalence relations $EQ(V)$ that punishes any deviation from the previously given conditions.

$$
\begin{aligned}
m(\sim) = (p_1 \cdot & |V/\sim| \\
+ p_2 \cdot & |\{v \in V/\sim | \exists x, y \in \bar{v} : op(x) \neq op(y)\}| \\
+ p_3 \cdot & \left|\{\bar{v} \in V/\sim | \exists y \in \bar{v}, x, n : x \xrightarrow{n} y \land \neg \bar{x} \xrightarrow{n} \bar{y})\}\right| \\
+ p_4 \cdot & |\{\bar{v} \in V/\sim | \text{ third conditions not satisfied}\}|)^{-p_5}
\end{aligned}
$$

The first parameter controls how much the equivalence relation is encouraged to merge together vertices into larger equivalence classes. The next three parameters control how strictly each property is enforced. The fifth parameter is simply used to change the distribution of the resulting scores without changing their relative order.

By trial and error we assigned the following values to each parameter: $p_1 = 1.0$, $p_2 = 0.5$, $p_3 = 0.5$, $p_4 = 0.5$.

2) Optimizing the Metric: Having defined the metric, we now need to find the optimal solution, i.e. the partition of the graph that maximizes the metric. We use a simple evolutionary algorithm over 1000 generations to maximize m.

C. Generating Constraints from a Matching

Given an appropriate equivalence relation, we can emit a constraint program in a straightforward fashion. To do this, we firstly generate a new graph $G/\sim := (V/\sim, E/\sim)$. Here, the set $E/\sim \subset \mathbb{N} \times V/\sim \times V/\sim$ is defined by the following property, where $(n, \bar{a}, \bar{b}) \in E/\sim$ iff $\bar{a} \xrightarrow{n} \bar{b}$.

We then define a threshold d and remove all vertices of the graph $G/\sim$ that contain fewer than d elements. This results in the removal of specific quirks of the individual synthesized programs, leaving the essential algorithmic skeleton intact. For the value of d, we choose the number of merged programs.

1) Example: Consider the example in Figure 5. On the left, we can see two simplified pieces of LLVM IR code from different loops. The first is from the body of a loop that computes the dot product of two vectors given as pointers, the second is from the innermost loop of a naive matrix multiplication kernel. They are matched together with the introduced graph matching algorithm, resulting in the clustering that is displayed in the middle column of the figure.

After discarding all equivalence classes with fewer than two elements, we generate constraints mechanically. Features that were specific to one program are discarded, notably the instructions %dot, %acc, %12 and the store instruction as seen in Figure 5. We can now use the generated constraints to find equivalent code sections in other source code [13].

D. Suggesting Replacements

Given a set of generated constraints as described above, the final step in our process is to search through user code for matching instances. The SMT-based search algorithm is described in detail by previous work [13], [15]; we treat this process as a black box.

We first ensure that the generalized constraints include a mapping for each function argument (if this is not the

Name	Kernels	Acceleration	LoC
NWChem	Dense	BLAS	1.2M
Abinit	Dense	BLAS,CUDA	900k
Pathsample	Sparse	Handcoded SpMV	40k
Darknet	Neural Network	CUDA	27k
Parboil	Linear Algebra	Handcoded MxM	187

TABLE I: Application source code used for evaluation.

Library	Platform	Kernels
Intel MKL	Intel CPU	Dense Linear Algebra
cuBLAS	Nvidia GPU	Linear Algebra
cuDNN	Nvidia GPU	Neural Networks
cuSparse	Nvidia GPU	Sparse Linear Algebra
CLBlast	OpenCL Devices	Dense Linear Algebra

TABLE II: Optimized libraries selected for evaluation.

case, we discard the constraints). From a search result, we can then identify which matched value corresponds to each function argument. Previous work targeting precisely known library functions [13] makes an automated replacement at this point by operating at the compiler IR level. However, as we require the developer to "sign off" on a replacement, we map the IR values back to source locations. We do this by exploiting Clang SSA value naming with the `-fno-discard-value-names` flag, as well as line and character debug information. This allows us to identify replacement opportunities.

VI. Setup

A. Applications

We selected compute intense applications that are likely to benefit from acceleration libraries. These are shown in Table I. Three of these are scientific applications: NWChem, Abinit and Pathsample, each representing a large code base. hese are significant applications; Abinit alone has been cited more than 6,000 times since 2002.

Abinit must be linked to a BLAS implementation installed somewhere on the target system, while NWChem uses an internal library and Pathsample implements a small set of required operations by hand (including some sparse methods). For each scientific application, we evaluated two different standard data sets which correspond to different chemistry simulations. These data sets were selected to exercise different portions of the underlying code base.

Neural networks also make large use of acceleration libraries. We examined Darknet [4], a widely used, open-source deep learning framework. It has been used recently to implement a number of highly cited, state-of-the-art models [16], [17], [18]. The framework consists of two distinct implementations (in C and CUDA) We evaluated three well known ImageNet models implemented using Darknet.

Finally, we examined a well known benchmark program from Parboil, SGEMM, which contains a hand-coded matrix-multiplication. As this is dominated by one linear algebra routine, it gives an upper bound on the typical performance achievable by our approach.

B. Libraries

The libraries (see Table II) fall broadly into two categories: those that are optimized for a particular CPU architecture (Intel MKL) to achieve performance, and those that use the GPU (CUDA libraries, CLBlast). There are a number of different

CUDA libraries that can be run on NVidia GPUs; for brevity we refer to these together as a single collection.

C. Platform

We targeted an Intel Xeon E5-2620 processor with 24 cores, 16GB of RAM and an NVidia Tesla K20 GPU. Applications were compiled at `-O3`. For the cross platform evaluation we targeted a 12 core AMD A10-7859k with an integrated AMD Radeon R7 iGPU.

D. Methodology

We ran each application from its "out of the box" configuration on the Intel platform to give a performance baseline.

For Pathsample the baseline code is sequential, handwritten Fortran with no library calls. This is also the case for NWChem which contains sequential C (which has hand inlined specimen BLAS implementations) In the case of Abinit, the baseline links to standard BLAS libraries. In the case of Darknet, there are 2 baselines available and we use the default sequential C baseline. We evaluated our performance when targeting both Intel MKL 2019 and a range of CUDA 8.0 based libraries.

To evaluate the impact of moving to a new platform, we focused on Darknet and evaluated our approach on an AMD platform that does not support CUDA. Instead it targets the CLBlast library.

VII. Results

We first evaluate the performance of our approach across applications, libraries and platforms. Next, we examine the number of library calls and candidate matches for API migration. This is followed by an evaluation of the execution time needed to synthesize equivalent programs for each library. Finally, we evaluate the accuracy of our graph matching algorithm and discuss the potential for unsound behavior to arise when using these tools.

A. Overall Results

A summary of our performance results is shown in Figure 6. On scientific applications, we found that the best implementation for each one achieved speedups of between 1.2 and $2.7\times$. This is the end to end performance of each application rather than just isolated kernels. In the case of Pathsample, the NGT configuration spends less time in sparse matrix operations than the PFold configuration. Amdahl's law means that inevitably PFold will benefit more from acceleration. MKL outperforms
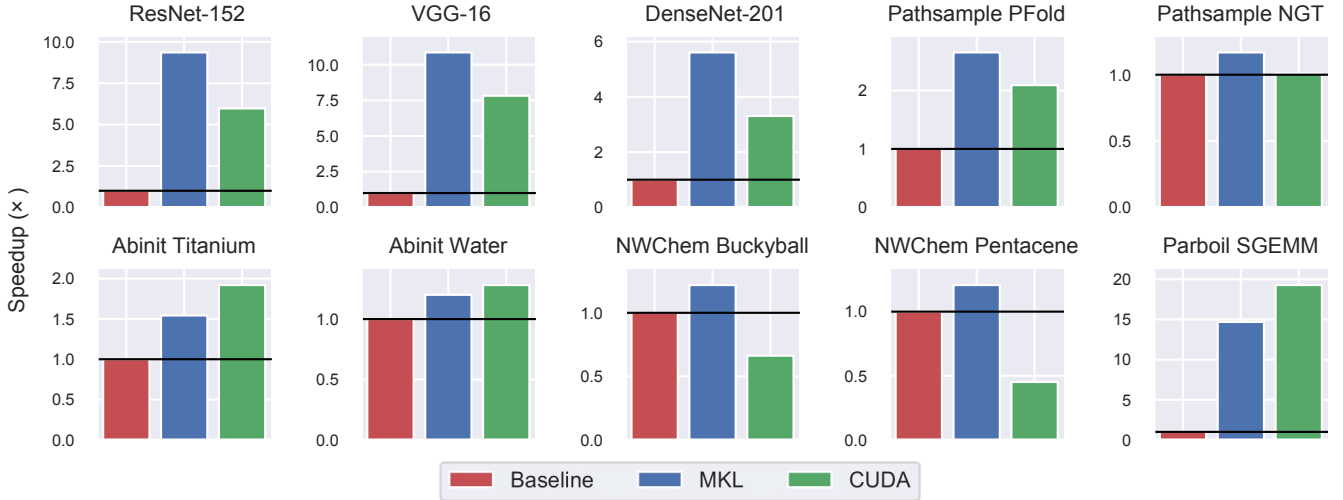
Fig. 6: Performance achievable by adopting code replacements suggested by our tools, for both Intel MKL and Nvidia CUDA libraries across the set of benchmarks listed.

the Nvidia libraries by a small margin in both cases. If we only used Nvidia libraries, there would be speedup available in the case of NGT.

This pattern continues with NWChem where MKL significantly outperforms the Nvidia libraries. Modest speedups are available for both configurations with an end-end speedup of $1.2\times$ Abinit shows a different behavior, where the Nvidia libraries outperform MKL, giving 1.2 to $1.9\times$ speedup. This is possibly due to the increased array sizes where the benefits of acceleration outweigh communication overhead. Unlike NWChem, both acceleration libraries improve performance.

We see more significant improvement for the DNNs as the amount of time spent in accelerator code sections is much greater. Improvements range from $5.5\times$ for the smaller DenseNet-201 to $11\times$ for the largest network: VGG-16. Like Pathsample and NWChem, all the the DNNs achieve the greatest performance with MKL, though Nvidia libraries still give improvements: $3.2\times$ to $7.7\times$. The impact of Amdahl's law can be clearly seen for Parboil SGEMM. Here there is just one kernel that can be readily accelerated. It achieves $15\times$ to $19\times$ speedups and provides a best case example.

B. Porting to New Hardware

Within Darknet, the use of optimized GPU libraries is built into the code: CUDA and CPU implementations are mixed together using preprocessor directives and the build system. As CUDA is not available on AMD GPU platforms, porting Darknet to such a platform means targeting OpenCL based libraries such as the CLBlast library [19], the results of which are shown in Figure 7.

We compared the performance of "out-of-the-box" Darknet against a handwritten OpenMP version [20], and our approach. The results of this comparison are shown in Figure 7. On all three networks, our approach outperforms the OpenMP


Fig. 7: Performance results for neural network inference on an AMD device with no CUDA support.

implementation which represents the best readily-available CPU performance on an AMD processor. We achieve speedups from $2.4\times$ on DenseNet-201 to $9\times$ on VGG-16. DenseNet-201 performs smaller matrix multiplications than the other networks, and so benefits less from GPU execution. Our results show that our approach allows for programmers to port applications to other platforms, without having to support multiple code bases for each possible implementation.

C. Library API usage

Table III shows the number of library call sites we detected in the original applications. For simplicity. we group functions that perform the same abstract computation together. For example, `cublas_sgemm`, `cblas_sgemm` and `clblast::Gemm<float>` are all considered together in the GEMM group.

Some of the applications we examined make extensive use of library functions. For example, Abinit links against an installed standard BLAS library, and so all the instances we detect in its code are from inlined library calls. Other applications bundle their own implementations; our approach detects this code rather than the corresponding call sites which results in a smaller overall number of matches. The true

		SPMV	GEMM	GEMV	GER	AXPY	AXPBY	SCAL	COPY	DOT	SOFTMAX	RELU
Abinit	**P**		180 (180)	47 (47)		21 (21)	2 (2)	20 (20)	70 (70)			
	TP		0/0/180/180	0/0/47/47		21/21/21/21	0/2/2/2	20/20/20/20	70/70/70/70			
	FP											
	FN		180/180/0/0	47/47/0/0			2/0/0/0					
Pathsample	**P**	2 (0)	1 (0)	1 (0)	3 (0)	7 (0)		13 (0)	5 (0)	1 (0)		
	TP	0/0/2/2	0/0/1/1	0/0/1/1	3/3/3/3	7/7/7/7		13/13/13/13	5/5/5/5	1/1/1/1		
	FP											
	FN	2/2/0/0	1/1/0/0	1/1/0/0								
NWChem	**P**	2 (0)	2 (0)	2 (0)	2 (0)	2 (0)	27 (0)	2 (0)	2 (0)	2 (0)		
	TP	0/0/2/2	0/0/2/2	0/0/2/2	0/2/2/2	0/2/2/2	0/27/27/27	0/2/2/2	0/2/2/2	0/2/2/2		
	FP							0/2/2/0	0/5/5/0			
	FN	2/2/0/0	2/2/0/0	2/2/0/0	2/0/0/0	2/0/0/0	27/0/0/0	2/0/0/0	2/0/0/0	2/0/0/0		
Darknet	**P**		2 (1)	1 (0)		1 (0)		1 (0)	1 (0)	1 (0)	1 (0)	1 (0)
	TP		0/0/2/2	0/0/1/1		0/1/1/1		0/1/1/1	0/1/1/1	0/1/1/1	0/0/0/0	0/1/1/1
	FP							0/3/3/0	0/2/2/0	0/1/1/0		
	FN		2/2/0/0	1/1/0/0		1/0/0/0		1/0/0/0	1/0/0/0	1/0/0/0	1/1/1/1	1/0/0/0
Parboil	**P**		1 (0)									
	TP		0/1/1/1									
	FP											
	FN		1/0/0/0									

TABLE III: Instances of each function category discovered across the different applications evaluated. The first row for each application gives the total number of potential matches we identified by hand-examination (positives, **P**), with the number of these corresponding to inlined library calls given in parentheses. Subsequent rows give the number of correctly identified opportunities (true positives, **TP**), incorrect matches (false positives, **FP**) and missed opportunities (false negatives, **FN**). Results are quoted as $x/y/z/w$ for the four different versions of our discovery algorithm: no generalization, generalization, generalization with nested loop corrections, and false positive testing. See Section VII-E for details.

positive figures (**P**) in Table III show both the total number of potential matches and the number that come from inlined library calls.

D. Synthesis

The time taken for our synthesizer to correctly synthesize each library program is acceptable for our usage model: every example could be synthesized in under 4 hours on a desktop-class machine, with examples that use shorter instruction sequences taking far less time. Synthesis time was not a primary goal of our work—learning the behavior of a function is a one-off task. If synthesis time is a bottleneck, there are many existing approaches to improve performance [21], [22], but this is beyond the scope of this paper.

E. Matching

Table III shows the results obtained when searching for code satisfying our generated constraints. We tested four different versions of the constraints: those generated from a single example constraint-based program, generalized versions from multiple programs, generalized with a post-processing step, and finally with dynamic testing of replacements.

The constraints generated from a single program fail to discover many examples. Only simple, inlined library calls are consistently matched by these constraints (Abinit in Table III, as the inlined code is identical to the code from which constraints are generated.

We then applied our graph matching algorithm to generalize constraints. These constraints are more successful; many instances that were not previously matched now are (e.g. in Darknet). Some instances such as GEMM and SPMV were not discovered by the generalized constraints. We discovered

this was due to a consistent difference between Clang's code generator and the synthesizer for nested loops—a mechanical post-processing step fixed these constraints, allowing these examples to be detected properly (GEMM, GEMV, SPMV columns in Table III).

Although these constraints generalized well, some false positive matches occur due to over-generalization (e.g. for SCAL, COPY in NWChem and Darknet). We found that this was due to missing data dependencies in code that interleaved another task with the learned function. To address this, we performed dynamic testing of suggested replacements using IO examples, eliminating all false positives we observed.

The only example not to be discovered in any of our test codes was SOFTMAX: it was implemented in the code we examined using a common numerical trick where the input data is shifted uniformly by its maximum value. The synthesizer is not able to learn this approach. Fortunately, it was not a significant contribution to execution time in the programs examined.

F. Soundness

There are a number of ways in which unsound behavior can arise when using our synthesis, generalization and replacement suggestion tools:

- Random IO examples may not capture the full range of a function's behavior. While this is a theoretical limitation of our synthesizer, in practice we have not encountered any function that suffers from this problem.
- The synthesizer may fail to synthesize a library function at all: not all functions have behavior that can be captured by the fragments used in this paper. If this happens, the

function is ignored. Our technique demonstrates useful performance improvements despite not being able to learn every library function.

- False positives and negatives can occur when matching constraints. We found that our constraints generalized well to detect complex examples, and that false positives can be readily eliminated by dynamic testing.

While these sources of unsound behavior can and do affect our process in some cases, the actual effects are not critical to the practical application of our tools.

G. Scalability

A natural question to be asked of our tooling is how well it scales beyond the examples shown in our existing evaluation. For example, in the context of machine learning we may wish to synthesize the behavior of a batch normalization layer or pooling. We anticipate that our approach will scale well to problems such as these. Other, more different problem domains such as sorting algorithms (a topic of interest in the program synthesis literature [23]) could be synthesized by expanding the set of fragments.

VIII. RELATED WORK

Rejuvenating Code: Most approaches within the software engineering community use a form of user guided refactoring [24] to perform simple syntactic restructuring of application code, supported by standard IDEs such as Eclipse [25]. In [26], replaying of refactoring for changed APIs calls is presented. However, refactoring techniques assume complete knowledge of library behavior. Such approaches do not address matching existing user code to emerging library APIs [27].

Automatic Acceleration: Our work uses the IDL constraint language [13] as a way to express computational patterns shared by a set of synthesized programs, and we make use of the detection features of IDL to discover instances of user code that match these patterns. While we automatically generate constraint descriptions for library interfaces based on their synthesized behavior, [13] requires an expert compiler developer to write constraint descriptions by hand; they do not consider synthesizing examples to drive constraint generation nor the generalization of constraints.

Sketching: Our technique uses control flow fragments to express partial guesses of the structure of a possible solution. [28] introduces the idea of *sketching* to allow the programmer to express a partial solution to a synthesis problem. Recent work in this area deals with recursive tree transformations [29] and modularity of sketches [30].

Our work lies within the area of *sketch generation*, where the synthesis problem is split into two parts: the first where sketches are generated, and the second where they are instantiated to produce solutions. Our approach lies within this space. SCYTHE [31] uses this technique to generate SQL queries, and the LASY language [32] uses libraries of composable domain-specific functions to describe the space of possible solutions, but requires a carefully-chosen set of input-output examples to work effectively.

Synthesizing Imperative Programs: Program synthesis often deals with synthesis of highly composable functional programs, which often allows synthesized programs to be represented in a minimal normal form [33]. Our technique generates LLVM intermediate representation, which does not permit a 'normal form' in the same way. Other work in which imperative programs are synthesized often deals only with straight line code [21], [34], or treats imperative control flow as special-cased components [35]. More commonly targeted than LLVM IR is assembly code, especially for superoptimization: [36], [22], [37], [38]. Our combination of target representation and treatment of control flow is a novel one.

Synthesizing for Acceleration: Other work has used program synthesis as a mechanism by which programs can be automatically accelerated. For example, recent work uses program synthesis to generate parallel versions of sequential code, with a focus on numerical array-processing algorithms with single-pass control flow [39], [40]. The space of programs tackled however is highly restricted. Helium [41] uses syntax-guided synthesis to synthesize programs in the Halide [42] image processing DSL. This approach has explicit knowledge of stencils hardwired into the synthesis phase and cannot be applied to unknown libraries.

Similarly, in [43], loop verification conditions are extracted from Fortran programs using inductive program synthesis techniques. These conditions can then be translated mechanically to Halide. Our approach to program synthesis is broader in application at the expense of formal verification.

Type-Directed Synthesis Our annotated type signatures are similar in spirit to *type-directed synthesis*. MYTH [44] uses type signatures alongside examples to synthesize recursive functional programs. More similar to our approach is the idea of using refinement types to guide the search process [45].

IX. CONCLUSION AND FUTURE WORK

Porting existing code to exploit accelerator libraries is a challenging problem for programmers. Understanding the behavior of existing and new libraries requires significant work on the programmer's part.

This paper presents two main contributions to help with this API evolution: a program synthesis technique that uses vendor-supplied type annotations to infer partial control flow structure for potential solutions, and a novel graph-matching based approach to finding a common description for a set of input programs. Using this approach we were able to achieve significant improvements to existing large scale code bases.

While our approach uses code normalization to aid matching, future work will focus on improving program synthesis to more closely match user code. This can be achieved by adding priors over the synthesis search space to bias construction. Currently the graph-matching algorithm introduces false positives that are eliminated with dynamic information. Future work will investigate the use of iteratively generating programs from generalized constraints and testing for equivalence against synthesized examples using SMT solvers, eliminating false positives.

REFERENCES

[1] A. Sedova, A. Tharrington, and B. Messer, "Portability in scientific computing: The molecular dynamics non-bonded forces calculation as a case study," 08 2018.

[2] D. Pflüger, M. Mehl, J. Valentin, F. Lindner, D. Pfander, S. Wagner, D. Graziotin, and Y. Wang, "The scalability-efficiency/maintainability-portability trade-off in simulation software engineering: Examples and a preliminary systematic literature review," in *Proceedings of the Fourth International Workshop on Software Engineering for HPC in Computational Science and Engineering*, ser. SE-HPCCSE '16. Piscataway, NJ, USA: IEEE Press, 2016, pp. 19–27. [Online]. Available: https://doi.org/10.1109/SE-HPCCSE.2016.8

[3] "PyTorch," https://pytorch.org/.

[4] J. Redmon, "Darknet: Open Source Neural Networks in C," 2013. [Online]. Available: http://pjreddie.com/darknet/

[5] "NVIDIA cuBLAS," https://developer.nvidia.com/cublas.

[6] "clMathLibraries clBLAS," https://github.com/clMathLibraries/clBLAS.

[7] "NVIDIA cuDNN," https://developer.nvidia.com/cudnn, 2018.

[8] "Intel® Compute Library for Deep Neural Networks (clDNN)," https://01.org/cldnn.

[9] W. Lee, K. Heo, R. Alur, and M. Naik, "Accelerating Search-based Program Synthesis Using Learned Probabilistic Models," in *Proceedings of the 39th ACM SIGPLAN Conference on Programming Language Design and Implementation*, ser. PLDI 2018. New York, NY, USA: ACM, 2018, pp. 436–449. [Online]. Available: http://doi.acm.org/10.1145/3192366.3192410

[10] M. Valiev, E. J. Bylaska, N. Govind, K. Kowalski, T. P. Straatsma, H. J. J. Van Dam, D. Wang, J. Nieplocha, E. Apra, T. L. Windus, and W. A. de Jong, "NWChem: A comprehensive and scalable open-source solution for large scale molecular simulations," *Computer Physics Communications*, vol. 181, no. 9, pp. 1477–1489, Sep. 2010. [Online]. Available: http://www.sciencedirect.com/science/article/pii/S0010465510001438

[11] "Intel® Math Kernel Library (MKL)," https://software.intel.com/mkl.

[12] S. Jha and S. A. Seshia, "A theory of formal synthesis via inductive learning," *Acta Informatica*, vol. 54, no. 7, pp. 693–726, Nov 2017. [Online]. Available: https://doi.org/10.1007/s00236-017-0294-5

[13] P. Ginsbach, T. Remmelg, M. Steuwer, B. Bodin, C. Dubach, and M. F. P. O'Boyle, "Automatic Matching of Legacy Code to Heterogeneous APIs: An Idiomatic Approach," in *Proceedings of the Twenty-Third International Conference on Architectural Support for Programming Languages and Operating Systems*, ser. ASPLOS '18. New York, NY, USA: ACM, 2018, pp. 139–153. [Online]. Available: http://doi.acm.org/10.1145/3173162.3173182

[14] K. Claessen and J. Hughes, "QuickCheck: A Lightweight Tool for Random Testing of Haskell Programs," in *Proceedings of the Fifth ACM SIGPLAN International Conference on Functional Programming*, ser. ICFP '00. New York, NY, USA: ACM, 2000, pp. 268–279. [Online]. Available: http://doi.acm.org/10.1145/351240.351266

[15] P. Ginsbach, L. Crawford, and M. F. P. O'Boyle, "CAnDL: A Domain Specific Language for Compiler Analysis," in *Proceedings of the 27th International Conference on Compiler Construction*, ser. CC 2018. New York, NY, USA: ACM, 2018, pp. 151–162. [Online]. Available: http://doi.acm.org/10.1145/3178372.3179515

[16] M. Rastegari, V. Ordonez, J. Redmon, and A. Farhadi, "XNOR-Net: ImageNet Classification Using Binary Convolutional Neural Networks," in *Computer Vision – ECCV 2016*, ser. Lecture Notes in Computer Science, B. Leibe, J. Matas, N. Sebe, and M. Welling, Eds. Springer International Publishing, 2016, pp. 525–542.

[17] J. Redmon, S. Divvala, R. Girshick, and A. Farhadi, "You Only Look Once: Unified, Real-Time Object Detection," in *2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, Jun. 2016, pp. 779–788.

[18] J. Redmon and A. Farhadi, "YOLO9000: Better, Faster, Stronger," in *2017 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, Jul. 2017, pp. 6517–6525.

[19] C. Nugteren, "CLBlast: A Tuned OpenCL BLAS Library," in *Proceedings of the International Workshop on OpenCL*, ser. IWOCL '18. New York, NY, USA: ACM, 2018, pp. 5:1–5:10. [Online]. Available: http://doi.acm.org/10.1145/3204919.3204924

[20] J. Turner, J. Cano, V. Radu, E. J. Crowley, M. O'Boyle, and A. Storkey, "Characterising across-stack optimisations for deep convolutional neural networks," in *2018 IEEE International Symposium on Workload Characterization (IISWC)*, Sep. 2018, pp. 101–110.

[21] S. Gulwani, S. Jha, A. Tiwari, and R. Venkatesan, "Synthesis of Loop-free Programs," in *Proceedings of the 32Nd ACM SIGPLAN Conference on Programming Language Design and Implementation*, ser. PLDI '11. New York, NY, USA: ACM, 2011, pp. 62–73. [Online]. Available: http://doi.acm.org/10.1145/1993498.1993506

[22] E. Schkufza, R. Sharma, and A. Aiken, "Stochastic Superoptimization," in *Proceedings of the Eighteenth International Conference on Architectural Support for Programming Languages and Operating Systems*, ser. ASPLOS '13. New York, NY, USA: ACM, 2013, pp. 305–316. [Online]. Available: http://doi.acm.org/10.1145/2451116.2451150

[23] C. D. Rosin, "Stepping Stones to Inductive Synthesis of Low-Level Looping Programs," *arXiv:1811.10665 [cs]*, Nov. 2018.

[24] D. Dig and R. Johnson, "The role of refactorings in api evolution," in *21st IEEE International Conference on Software Maintenance (ICSM'05)*. IEEE, 2005, pp. 389–398.

[25] G. C. Murphy, M. Kersten, and L. Findlater, "How are java software developers using the elipse ide?" *IEEE software*, vol. 23, no. 4, pp. 76–83, 2006.

[26] J. Henkel and A. Diwan, "Catchup! capturing and replaying refactorings to support api evolution," in *Proceedings. 27th International Conference on Software Engineering, 2005. ICSE 2005*. IEEE, 2005, pp. 274–283.

[27] A. Brito, L. Xavier, A. Hora, and M. T. Valente, "Apidiff: Detecting api breaking changes," in *2018 IEEE 25th International Conference on Software Analysis, Evolution and Reengineering (SANER)*. IEEE, 2018, pp. 507–511.

[28] A. Solar-Lezama, "The Sketching Approach to Program Synthesis," in *Programming Languages and Systems*, ser. Lecture Notes in Computer Science. Springer, Berlin, Heidelberg, Dec. 2009, pp. 4–13. [Online]. Available: https://link.springer.com/chapter/10.1007/978-3-642-10672-9_3

[29] J. P. Inala, N. Polikarpova, X. Qiu, B. S. Lerner, and A. Solar-Lezama, "Synthesis of Recursive ADT Transformations from Reusable Templates," in *Tools and Algorithms for the Construction and Analysis of Systems*, ser. Lecture Notes in Computer Science, A. Legay and T. Margaria, Eds. Springer Berlin Heidelberg, 2017, pp. 247–263.

[30] R. Singh, R. Singh, Z. Xu, R. Krosnick, and A. Solar-Lezama, "Modular Synthesis of Sketches Using Models," in *Verification, Model Checking, and Abstract Interpretation*, ser. Lecture Notes in Computer Science, K. L. McMillan and X. Rival, Eds. Springer Berlin Heidelberg, 2014, pp. 395–414.

[31] C. Wang, A. Cheung, and R. Bodik, "Synthesizing Highly Expressive SQL Queries from Input-output Examples," in *Proceedings of the 38th ACM SIGPLAN Conference on Programming Language Design and Implementation*, ser. PLDI 2017. New York, NY, USA: ACM, 2017, pp. 452–466. [Online]. Available: http://doi.acm.org/10.1145/3062341.3062365

[32] D. Perelman, S. Gulwani, D. Grossman, and P. Provost, "Test-driven Synthesis," in *Proceedings of the 35th ACM SIGPLAN Conference on Programming Language Design and Implementation*, ser. PLDI '14. New York, NY, USA: ACM, 2014, pp. 408–418. [Online]. Available: http://doi.acm.org/10.1145/2594291.2594297

[33] J. K. Feser, S. Chaudhuri, and I. Dillig, "Synthesizing Data Structure Transformations from Input-output Examples," in *Proceedings of the 36th ACM SIGPLAN Conference on Programming Language Design and Implementation*, ser. PLDI '15. New York, NY, USA: ACM, 2015, pp. 229–239. [Online]. Available: http://doi.acm.org/10.1145/2737924.2737977

[34] R. Sasnauskas, Y. Chen, P. Collingbourne, J. Ketema, G. Lup, J. Taneja, and J. Regehr, "Souper: A Synthesizing Superoptimizer," *arXiv:1711.04422 [cs]*, Nov. 2017. [Online]. Available: http://arxiv.org/abs/1711.04422

[35] S. Gulwani, "Automating String Processing in Spreadsheets Using Input-output Examples," in *Proceedings of the 38th Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*, ser. POPL '11. New York, NY, USA: ACM, 2011, pp. 317–330. [Online]. Available: http://doi.acm.org/10.1145/1926385.1926423

[36] V. Srinivasan and T. Reps, "Synthesis of Machine Code from Semantics," in *Proceedings of the 36th ACM SIGPLAN Conference on Programming Language Design and Implementation*, ser. PLDI '15. New York, NY, USA: ACM, 2015, pp. 596–607. [Online]. Available: http://doi.acm.org/10.1145/2737924.2737960

[37] S. Heule, E. Schkufza, R. Sharma, and A. Aiken, "Stratified Synthesis: Automatically Learning the x86-64 Instruction Set," in *Proceedings of the 37th ACM SIGPLAN Conference on Programming Language Design and Implementation*, ser. PLDI '16. New York, NY, USA: ACM, 2016, pp. 237–250. [Online]. Available: http://doi.acm.org/10.1145/2908080.2908121

[38] N. Hasabnis, R. Qiao, and R. Sekar, "Checking Correctness of Code Generator Architecture Specifications," in *Proceedings of the 13th Annual IEEE/ACM International Symposium on Code Generation and Optimization*, ser. CGO '15. Washington, DC, USA: IEEE Computer Society, 2015, pp. 167–178. [Online]. Available: http://dl.acm.org/citation.cfm?id=2738600.2738622

[39] A. Farzan and V. Nicolet, "Synthesis of Divide and Conquer Parallelism for Loops," in *Proceedings of the 38th ACM SIGPLAN Conference on Programming Language Design and Implementation*, ser. PLDI 2017. New York, NY, USA: ACM, 2017, pp. 540–555. [Online]. Available: http://doi.acm.org/10.1145/3062341.3062355

[40] G. Fedyukovich, M. B. S. Ahmad, and R. Bodik, "Gradual Synthesis for Static Parallelization of Single-pass Array-processing Programs," in *Proceedings of the 38th ACM SIGPLAN Conference on Programming Language Design and Implementation*, ser. PLDI 2017. New York, NY, USA: ACM, 2017, pp. 572–585. [Online]. Available: http://doi.acm.org/10.1145/3062341.3062382

[41] C. Mendis, J. Bosboom, K. Wu, S. Kamil, J. Ragan-Kelley, S. Paris, Q. Zhao, and S. Amarasinghe, "Helium: Lifting high-performance stencil kernels from stripped x86 binaries to halide dsl code," in *ACM SIGPLAN Conference on Programming Language Design and Implementation*, Portland, OR, June 2015. [Online]. Available: http://groups.csail.mit.edu/commit/papers/2015/mendis-pldi15-helium.pdf

[42] J. Ragan-Kelley, C. Barnes, A. Adams, S. Paris, F. Durand, and S. Amarasinghe, "Halide: A Language and Compiler for Optimizing Parallelism, Locality, and Recomputation in Image Processing Pipelines," in *Proceedings of the 34th ACM SIGPLAN Conference on Programming Language Design and Implementation*, ser. PLDI '13. New York, NY, USA: ACM, 2013, pp. 519–530. [Online]. Available: http://doi.acm.org/10.1145/2491956.2462176

[43] S. Kamil, A. Cheung, S. Itzhaky, and A. Solar-Lezama, "Verified Lifting of Stencil Computations," in *Proceedings of the 37th ACM SIGPLAN Conference on Programming Language Design and Implementation*, ser. PLDI '16. New York, NY, USA: ACM, 2016, pp. 711–726. [Online]. Available: http://doi.acm.org/10.1145/2908080.2908117

[44] P.-M. Osera and S. Zdancewic, "Type-and-example-directed Program Synthesis," in *Proceedings of the 36th ACM SIGPLAN Conference on Programming Language Design and Implementation*, ser. PLDI '15. New York, NY, USA: ACM, 2015, pp. 619–630. [Online]. Available: http://doi.acm.org/10.1145/2737924.2738007

[45] N. Polikarpova, I. Kuraj, and A. Solar-Lezama, "Program Synthesis from Polymorphic Refinement Types," in *Proceedings of the 37th ACM SIGPLAN Conference on Programming Language Design and Implementation*, ser. PLDI '16. New York, NY, USA: ACM, 2016, pp. 522–538. [Online]. Available: http://doi.acm.org/10.1145/2908080.2908093